

# Report for Homework1 of CS6650

## GitHub Repo

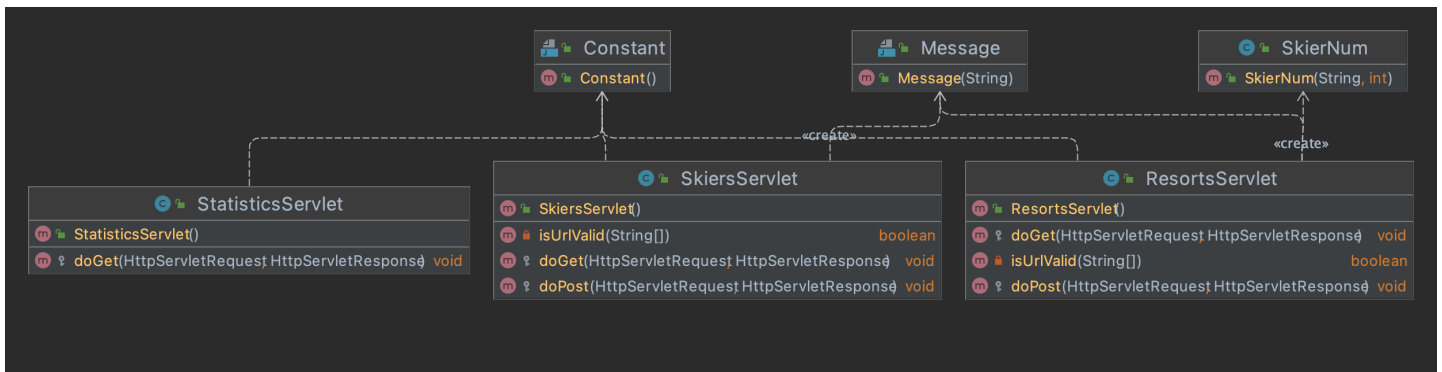
---

<https://github.com/zjdx1998/CS6650/tree/Homework1/Assignment1>

## Server

---

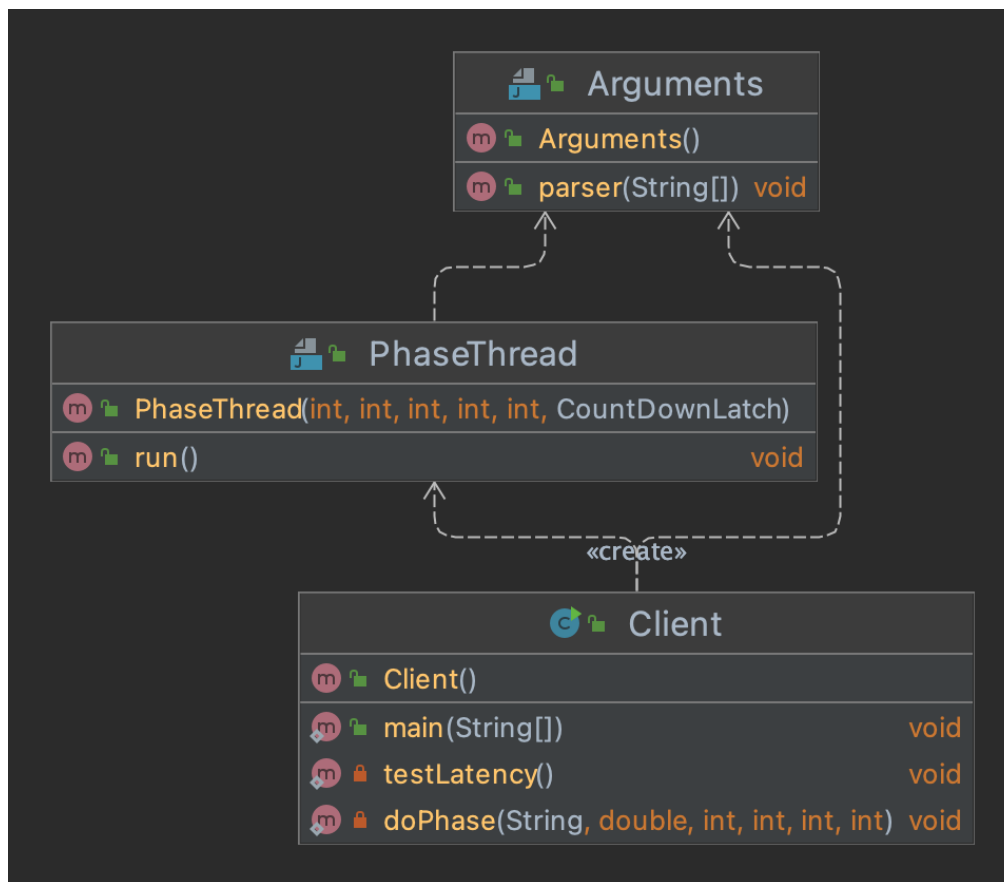
I implemented a full-completed server which can handle every request and return response as <https://app.swaggerhub.com/apis/cloud-perf/SkiDataAPI/1.16#/skiers/getSkierResortTotals> expected.



## Client Overview

---

# Design

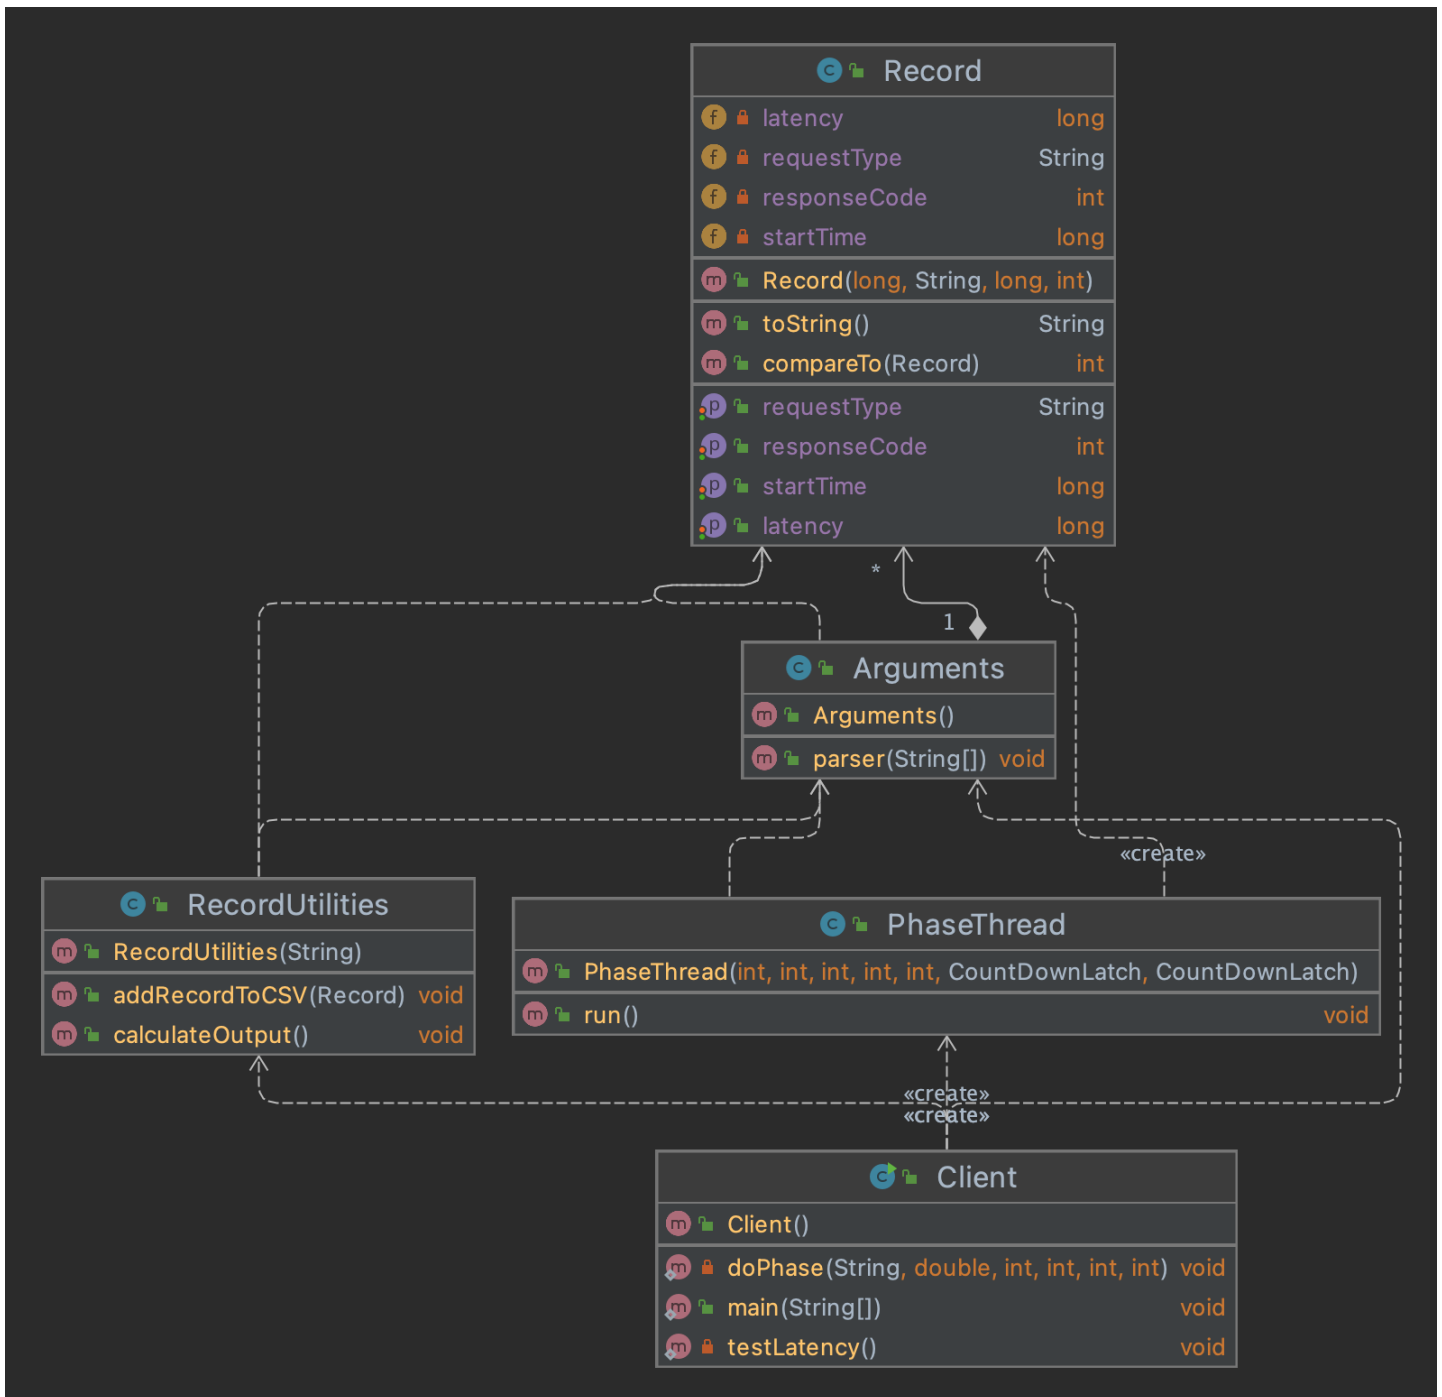


Client 1 is consist of three parts:

The first part is global arguments processing, where I utilized `args4j` to parse the arguments. I also stored the global variables like `successCount`, `failureCount` which represents the total number of successful and failed requests.

The second part is a runnable class `PhaseThread`, which accepts `int startID, int endID, int startTime, int endTime, int numOfReqs, CountDownLatch latch` as parameter, this is the per client request model.

The third part is main class `client` which has a `static void main` function. This class firstly load the arguments from `Argument` then create a thread pool with total numbers of three phases. All the statictics data are calculated here.



In addition to Client 1, Client 2 also contains a `csvUtiliizer` class to utilize the data.

# Little's Law Verification

```
success
Ready to run phases!
Phase1 is ready to start!
Phase1 should execute 8 threads with 500 requests each.
Phase1 has already completed 20.0% tasks
Phase2 is ready to start!
Phase2 should execute 32 threads with 375 requests each.
Phase2 has already completed 20.0% tasks
Phase3 is ready to start!
Phase3 should execute 3 threads with 1 requests each.
Phase3 has already completed 100.0% tasks
Number of Successful Requests Sent: 16003
Number of Unsuccessful Requests: 0
Total run time: 55928 (ms)
Total Throughput in requests per second: 286.1357459590983
```

```
/Users/jeromy/.sdkman/candidates/java/17-open/bin/java ...
Ready to test latency!
Total Duration is 116.917 with average latency about 58.4585
```

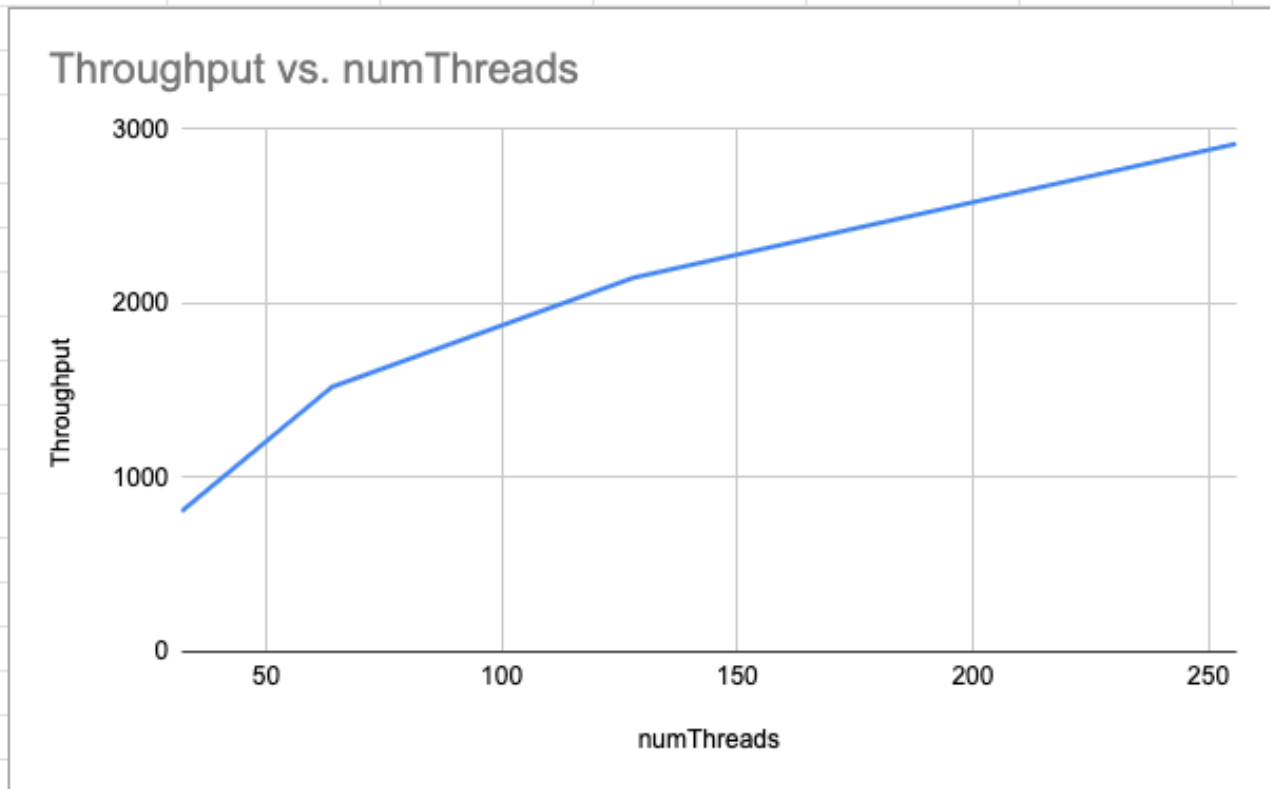
The above results are generated by `-nt 32 -ns 2000 -nl 40 -nr 10 -test-only -server ec2-54-149-212-65.us-west-2.compute.amazonaws.com` arguments with `-test-only` enabled or disabled. The test only data is calculated by test same number of threads

We can see the predicted throughput for test-only is  $\frac{32}{58.4585} \approx \frac{16003}{55928} = 0.286$ , especially consider three phases of clients using different number of threads.

## Client 1 Statistics

---

numThreads	32	64	128	256
Throughput	805.4274	1520.7671	2147.7313	2916.9799



We can see the throughput is linear related to the num of Threads.

```
Ready to run phases!
Phase1 is ready to start!
Phase1 should execute 8 threads with 5000 requests each.
Phase1 has already completed 20.0% tasks
Phase2 is ready to start!
Phase2 should execute 32 threads with 3750 requests each.
Phase2 has already completed 20.0% tasks
Phase3 is ready to start!
Phase3 should execute 3 threads with 1 requests each.
Phase3 has already completed 100.0% tasks
Number of Successful Requests Sent: 160003
Number of Unsuccessful Requests: 0
Total run time: 198656 (ms)
Total Throughput in requests per second: 805.4274726159794
```

```
Ready to run phases!
Phase1 is ready to start!
Phase1 should execute 16 threads with 2500 requests each.
Phase1 has already completed 20.0% tasks
Phase2 is ready to start!
Phase2 should execute 64 threads with 1875 requests each.
Phase2 has already completed 20.0% tasks
Phase3 is ready to start!
Phase3 should execute 6 threads with 1 requests each.
Phase3 has already completed 100.0% tasks
Number of Successful Requests Sent: 160006
Number of Unsuccessful Requests: 0
Total run time: 105214 (ms)
Total Throughput in requests per second: 1520.7671982815975

Ready to run phases!
Phase1 is ready to start!
Phase1 should execute 32 threads with 1250 requests each.
Phase1 has already completed 20.0% tasks
Phase2 is ready to start!
Phase2 should execute 128 threads with 937 requests each.
Phase2 has already completed 20.0% tasks
Phase3 is ready to start!
Phase3 should execute 12 threads with 1 requests each.
Phase3 has already completed 100.0% tasks
Number of Successful Requests Sent: 159948
Number of Unsuccessful Requests: 0
Total run time: 74473 (ms)
Total Throughput in requests per second: 2147.7313925852322
```

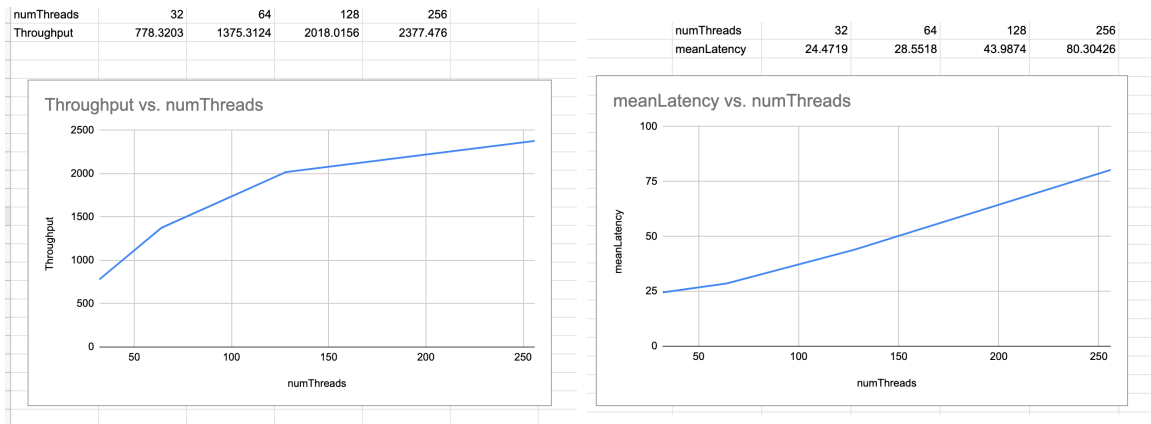
```

Ready to run phases!
Phase1 is ready to start!
Phase1 should execute 64 threads with 625 requests each.
Phase1 has already completed 20.0% tasks
Phase2 is ready to start!
Phase2 should execute 256 threads with 468 requests each.
Phase2 has already completed 20.0% tasks
Phase3 is ready to start!
Phase3 should execute 25 threads with 1 requests each.
Phase3 has already completed 100.0% tasks
Number of Successful Requests Sent: 159833
Number of Unsuccessful Requests: 0
Total run time: 54794 (ms)
Total Throughput in requests per second: 2916.9799613096325

```

The above charts and results are did by `-nt 32 -ns 2000 -nl 40 -nr 10 -server ec2-54-149-212-65.us-west-2.compute.amazonaws.com` with `nt` changes to `32/64/128/256` respectively.

## Clients 2 Statistics



We can see both Throughput and mean Latency are linearly related to the number of threads. Consider the default maximum restrictions of Tomcat in AWS EC2 is 200, the first chart contains a inflection point which represents this case.

Success

Ready to run phases!

Phase1 is ready to start!

Phase1 should execute 8 threads with 5000 requests each.

Phase1 has already completed 20.0% tasks

Phase2 is ready to start!

Phase2 should execute 32 threads with 3750 requests each.

Phase2 has already completed 20.0% tasks

Phase3 is ready to start!

Phase3 should execute 3 threads with 1 requests each.

Phase3 has already completed 100.0% tasks

Number of Successful Requests Sent: 157957

Number of Unsuccessful Requests: 0

Total run time: 202946 (ms)

Total Throughput in requests per second: 778.3203413715964

Mean response time: 24.47194473089844

Median response time: 20.0

Throughput: 40.863119420884985

99th response time: 94.0

min and max response time: min: 11.0 , max: 2575.0

Success

Ready to run phases!

Phase1 is ready to start!

Phase1 should execute 16 threads with 2500 requests each.

Phase1 has already completed 20.0% tasks

Phase2 is ready to start!

Phase2 should execute 64 threads with 1875 requests each.

Phase2 has already completed 20.0% tasks

Phase3 is ready to start!

Phase3 should execute 6 threads with 1 requests each.

Phase3 has already completed 100.0% tasks

Number of Successful Requests Sent: 156831

Number of Unsuccessful Requests: 0

Total run time: 114033 (ms)

Total Throughput in requests per second: 1375.3124095656522

Mean response time: 28.55185944885401

Median response time: 24.0

Throughput: 35.02398860541243

99th response time: 72.0

min and max response time: min: 12.0 , max: 3409.0



```

Success
Ready to run phases!
Phase1 is ready to start!
Phase1 should execute 32 threads with 1250 requests each.
Phase1 has already completed 20.0% tasks
Phase2 is ready to start!
Phase2 should execute 128 threads with 937 requests each.
Phase2 has already completed 20.0% tasks
Phase3 is ready to start!
Phase3 should execute 12 threads with 1 requests each.
Phase3 has already completed 100.0% tasks
Number of Successful Requests Sent: 156470
Number of Unsuccessful Requests: 0
Total run time: 77502 (ms)
Total Throughput in requests per second: 2018.9156408866868
Mean response time: 43.98747534327901
Median response time: 35.0
Throughput: 22.733743916785013
99th response time: 215.0
min and max response time: min: 11.0 , max: 3927.0

```


```

Success
Ready to run phases!
Phase1 is ready to start!
Phase1 should execute 64 threads with 625 requests each.
Phase1 has already completed 20.0% tasks
Phase2 is ready to start!
Phase2 should execute 256 threads with 468 requests each.
Phase2 has already completed 20.0% tasks
Phase3 is ready to start!
Phase3 should execute 25 threads with 1 requests each.
Phase3 has already completed 100.0% tasks
Number of Successful Requests Sent: 152672
Number of Unsuccessful Requests: 0
Total run time: 64216 (ms)
Total Throughput in requests per second: 2377.4760184377724
Mean response time: 80.30426003254328
Median response time: 52.0
Throughput: 12.45263949377966
99th response time: 607.0
min and max response time: min: 12.0 , max: 4884.0

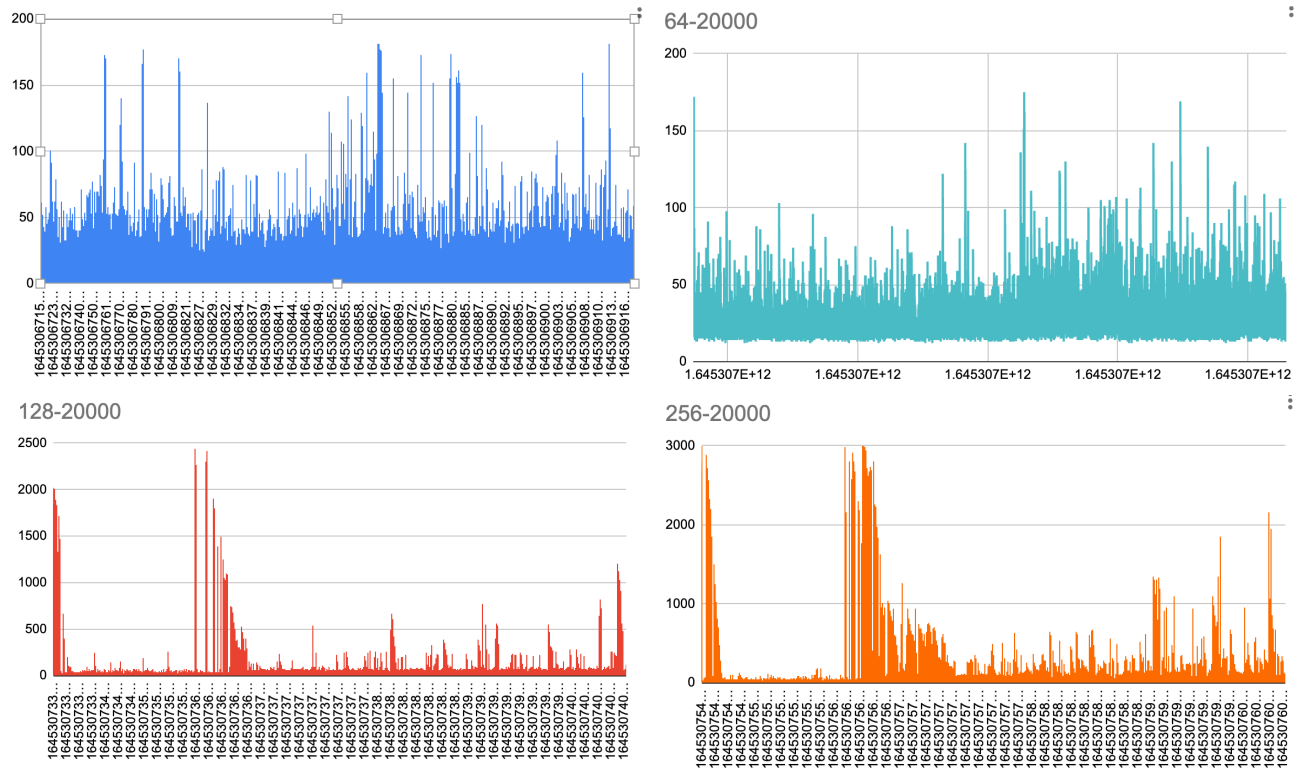
```

## Bonus

### Latency(50%, 99%)



The above chart is plotted for `numThreads=128`, `numSkiers=20000`, we can clearly see the p50 and p99, but it's not easy to figure out how it's distributed.



Now we know the latency distribution is very interesting which matches well with the actual request layout. This is exactly plotted by average response time for certain time interval bucket.

All the raw data can be found at [Clients/res/records](#).